

# ComputingTheEuclideanDistance

August 8, 2023

## 1 Computing the Euclidean Distance

```
[6]: import pandas as pd
import numpy as np
import os
import math
import random
import matplotlib.pyplot as plt
```

### 1.0.1 Load a Data Set and Save it as a Pandas DataFrame

We will work with a new data set called "cell2cell." This data set is used to analyze cellular telephone customers and can be used to predict whether a customer will remain with their current telecom service or leave to another.

```
[7]: filename = os.path.join(os.getcwd(), "data", "cell2cell.csv")
df = pd.read_csv(filename, header=0)
```

### 1.0.2 Inspect the Data

```
[8]: df.head()
```

```
[8]: CustomerID Churn ServiceArea ChildrenInHH HandsetRefurbished \
0      3000002   True  SEAPOR503          False             False
1      3000010   True  PITHOM412           True             False
2      3000014  False  MILMIL414           True             False
3      3000022  False  PITHOM412          False             False
4      3000026   True  OKCTUL918          False             False

HandsetWebCapable TruckOwner RVOwner HomeownershipKnown \
0                True      False    False              True
1                False      False    False              True
2                False      False    False             False
3                True      False    False              True
4                False      False    False              True

BuysViaMailOrder ... HandsetModels CurrentEquipmentDays  AgeHH1 \
```

0	True	...	0.487071	-0.077013	1.387766
1	True	...	-0.616775	3.019920	0.392039
2	False	...	-0.616775	3.019920	-0.241605
3	True	...	2.694763	0.305179	-0.060564
4	True	...	1.590917	1.857585	0.663601

	AgeHH2	RetentionCalls	RetentionOffersAccepted	\
0	-0.883541	4.662897	-0.1283	
1	0.871495	-0.180167	-0.1283	
2	0.202910	-0.180167	-0.1283	
3	-0.883541	-0.180167	-0.1283	
4	1.372934	-0.180167	-0.1283	

	ReferralsMadeBySubscriber	IncomeGroup	AdjustmentsToCreditRating	\
0	-0.169283	-0.103411	-0.140707	
1	-0.169283	0.215243	-0.140707	
2	-0.169283	0.533896	-0.140707	
3	-0.169283	0.533896	-0.140707	
4	-0.169283	1.489856	2.469282	

	HandsetPrice
0	-0.864858
1	-0.864858
2	-0.368174
3	-1.195980
4	-1.195980

[5 rows x 58 columns]

[9]: df.dtypes

```
[9]: CustomerID          int64
      Churn              bool
      ServiceArea        object
      ChildrenInHH        bool
      HandsetRefurbished  bool
      HandsetWebCapable   bool
      TruckOwner          bool
      RVOwner             bool
      HomeownershipKnown  bool
      BuysViaMailOrder     bool
      RespondsToMailOffers bool
      OptOutMailings       bool
      NonUSTravel         bool
      OwnsComputer         bool
      HasCreditCard       bool
      NewCellphoneUser     bool
      NotNewCellphoneUser  bool
```

OwnsMotorcycle	bool
MadeCallToRetentionTeam	bool
CreditRating	object
PrizmCode	object
Occupation	object
Married	object
MonthlyRevenue	float64
MonthlyMinutes	float64
TotalRecurringCharge	float64
DirectorAssistedCalls	float64
OverageMinutes	float64
RoamingCalls	float64
PercChangeMinutes	float64
PercChangeRevenues	float64
DroppedCalls	float64
BlockedCalls	float64
UnansweredCalls	float64
CustomerCareCalls	float64
ThreewayCalls	float64
ReceivedCalls	float64
OutboundCalls	float64
InboundCalls	float64
PeakCallsInOut	float64
OffPeakCallsInOut	float64
DroppedBlockedCalls	float64
CallForwardingCalls	float64
CallWaitingCalls	float64
MonthsInService	float64
UniqueSubs	float64
ActiveSubs	float64
Handsets	float64
HandsetModels	float64
CurrentEquipmentDays	float64
AgeHH1	float64
AgeHH2	float64
RetentionCalls	float64
RetentionOffersAccepted	float64
ReferralsMadeBySubscriber	float64
IncomeGroup	float64
AdjustmentsToCreditRating	float64
HandsetPrice	float64
dtype:	object

```
[10]: df.shape
```

```
[10]: (51047, 58)
```

## 1.1 Euclidean Distance

**KNN** k-Nearest Neighbors (KNN) is an instance-based learning algorithm. To make a classification for a given unlabeled example  $A$ , we search the training data for the  $k$  nearest neighbors, as defined by some distance metric  $d(A, B)$  in which  $B$  represents another example. We choose the most common label among the nearest neighbor examples to be our prediction (label) for the unlabeled example.

The most commonly used distance metric for KNN is the Euclidean distance.

**Euclidean Distance** For two  $n$ -dimensional, real-valued vectors  $A, B \in \mathbb{R}^n$ , the Euclidean distance  $eud$  is defined as:

$$eud(A, B) = \sqrt{\sum_{i=1}^n (B_i - A_i)^2}$$

Euclidean distance finds the distance between two vectors of the same length. In this formula,  $A_i$  is the  $i$ th coordinate of vector  $A$ , and  $B_i$  is the  $i$ th coordinate of vector  $B$ .

Let's relate this to a dataset. Let's think of the vectors  $A$  and  $B$  as being two examples (rows) in a dataset.

Let  $A = \langle x_1^a, \dots, x_n^a \rangle$  be a  $n$ -dimensional vector ( $x_i^a$  is the  $i$ th feature in example  $A$  and  $n$  is the total number of features).

Then for two vectors (examples)  $A$  and  $B$  the Euclidean distance is defined as:

$$eud(A, B) = \sqrt{(x_1^b - x_1^a)^2 + (x_2^b - x_2^a)^2 + \dots + (x_n^b - x_n^a)^2} = \sqrt{\sum_{i=1}^n (x_i^b - x_i^a)^2}$$

To visualize KNN, you can picture plotting the examples (also called data points) in our dataset and finding the distance between them. Let's create a visualization to see how we plot examples and find the distance between each example.

To easily visualize this, let's plot two examples from DataFrame `df`. Note that each example contains many features, but to make this visualization even simpler, we will work with two dimensions (that is, two features).

Euclidean distance is best used to calculate the distance between vectors containing numerical values. Therefore, we will choose two features that have numerical values.

Let us use row 0 and row 4 in DataFrame `df` and focus on features `HandsetModels` and `AgeHH1`. Run the code below to examine the two examples we will be plotting.

```
[11]: display(df.loc[[0,4], ['HandsetModels', 'AgeHH1']])
```

	HandsetModels	AgeHH1
0	0.487071	1.387766
4	1.590917	0.663601

Each example (row) can be viewed as a vector:

example 1: (0.487071, 1.387766)

example 2: (1.590917, 0.663601)

You will use the Euclidean distance formula to find the distance between these two vectors. First, let's plot these vectors. Run the code cell below to generate a plot. Examine the resulting plot.


```
[12]: # example 1 (row 0):
vector_A = [df.loc[0]['HandsetModels'], df.loc[0]['AgeHH1']]

# example 2 (row 4):
vector_B = [df.loc[4]['HandsetModels'], df.loc[4]['AgeHH1']]

plt.scatter(vector_A[0],vector_A[1] ,c='b',label='example 1')
plt.scatter(vector_B[0],vector_B[1], c='r', label='example 2')
plt.plot([vector_A[0],vector_B[0]], [vector_A[1],vector_B[1]], c='g',
         linestyle='dashed', label='distance')

plt.xlim([0, 3])
plt.ylim([0, 3])
plt.xlabel('HandsetModels')
plt.ylabel('AgeHH1')

plt.legend(loc='upper right');
plt.show()
```



You will use the Euclidean distance formula to find the distance between these two vectors. Use the Euclidean distance formula to calculate the distance between vector\_A and vector\_B by hand and save the result to variable `euc_distance`. For simplicity, use the following rounded vector values in your calculation:  
vector\_A: (0.5, 1.4)

```
vector_B: (1.6, 0.7)
```

### 1.1.1 Graded Cell

The cell below will be graded. Remove the line "raise NotImplementedError()" before writing your code.

```
[13]: # YOUR CODE HERE
euc_distance=1.303
```

### 1.1.2 Self-Check

Run the cell below to test the correctness of your code above before submitting for grading. Do not add code or delete code in the cell.

```
[14]: # Run this self-test cell to check your code;
# do not add code or delete code in this cell
from jn import testEuc

try:
    p, err = testEuc(euc_distance)
    print(err)
except Exception as e:
    print("Error!\n" + str(e))
```

Correct!

## 1.2 Step 1: Filter Numerical Features

We will now compute the Euclidean distance between two rows in DataFrame df, using all of their numerical feature values. Let us create a new DataFrame that contains only the numerically valued columns of the original df DataFrame.

```
[15]: df_numerical = df.select_dtypes(include=['int64', 'float64'])

print(df_numerical.shape)
df_numerical.head()
```

```
(51047, 36)
```

```
[15]: CustomerID  MonthlyRevenue  MonthlyMinutes  TotalRecurringCharge  \
0      3000002      -0.782676      -0.578738      -1.041153
1      3000010      -0.940180      -0.973177      -1.250809
2      3000014      -0.468118      -0.976952      -0.370255
3      3000022       0.526784       1.484048       1.181196
4      3000026      -0.936810      -0.992050      -1.250809

DirectorAssistedCalls  OverageMinutes  RoamingCalls  PercChangeMinutes  \
0      -0.289532      -0.414422      -0.125914      -0.564836
```

1	-0.401714	-0.414422	-0.125914	0.029311
2	-0.401714	-0.414422	-0.125914	0.037077
3	0.154708	-0.414422	-0.125914	0.654524
4	-0.401714	-0.414422	-0.125914	0.044844

	PercChangeRevenues	DroppedCalls	...	HandsetModels	CurrentEquipmentDays	\
0	-0.449987	-0.587303	...	0.487071	-0.077013	
1	0.030120	-0.631532	...	-0.616775	3.019920	
2	0.030120	-0.664703	...	-0.616775	3.019920	
3	0.234797	4.012499	...	2.694763	0.305179	
4	0.025066	-0.664703	...	1.590917	1.857585	

	AgeHH1	AgeHH2	RetentionCalls	RetentionOffersAccepted	\
0	1.387766	-0.883541	4.662897	-0.1283	
1	0.392039	0.871495	-0.180167	-0.1283	
2	-0.241605	0.202910	-0.180167	-0.1283	
3	-0.060564	-0.883541	-0.180167	-0.1283	
4	0.663601	1.372934	-0.180167	-0.1283	

	ReferralsMadeBySubscriber	IncomeGroup	AdjustmentsToCreditRating	\
0	-0.169283	-0.103411	-0.140707	
1	-0.169283	0.215243	-0.140707	
2	-0.169283	0.533896	-0.140707	
3	-0.169283	0.533896	-0.140707	
4	-0.169283	1.489856	2.469282	

	HandsetPrice
0	-0.864858
1	-0.864858
2	-0.368174
3	-1.195980
4	-1.195980

[5 rows x 36 columns]

We will exclude the CustomerID column, since it contains the customer ID and is not a feature that we want to consider.

```
[16]: df_numerical = df_numerical.drop(columns=['CustomerID'])
df_numerical.head()
```

```
[16]: MonthlyRevenue  MonthlyMinutes  TotalRecurringCharge  \
0      -0.782676      -0.578738      -1.041153
1      -0.940180      -0.973177      -1.250809
2      -0.468118      -0.976952      -0.370255
3       0.526784       1.484048       1.181196
4      -0.936810      -0.992050      -1.250809
```

	DirectorAssistedCalls	OverageMinutes	RoamingCalls	PercChangeMinutes	\
--	-----------------------	----------------	--------------	-------------------	---

0	-0.289532	-0.414422	-0.125914	-0.564836
1	-0.401714	-0.414422	-0.125914	0.029311
2	-0.401714	-0.414422	-0.125914	0.037077
3	0.154708	-0.414422	-0.125914	0.654524
4	-0.401714	-0.414422	-0.125914	0.044844

	PercChangeRevenues	DroppedCalls	BlockedCalls	...	HandsetModels \
0	-0.449987	-0.587303	-0.309284	...	0.487071
1	0.030120	-0.631532	-0.373230	...	-0.616775
2	0.030120	-0.664703	-0.373230	...	-0.616775
3	0.234797	4.012499	0.330172	...	2.694763
4	0.025066	-0.664703	-0.373230	...	1.590917

	CurrentEquipmentDays	AgeHH1	AgeHH2	RetentionCalls \
0	-0.077013	1.387766	-0.883541	4.662897
1	3.019920	0.392039	0.871495	-0.180167
2	3.019920	-0.241605	0.202910	-0.180167
3	0.305179	-0.060564	-0.883541	-0.180167
4	1.857585	0.663601	1.372934	-0.180167

	RetentionOffersAccepted	ReferralsMadeBySubscriber	IncomeGroup \
0	-0.1283	-0.169283	-0.103411
1	-0.1283	-0.169283	0.215243
2	-0.1283	-0.169283	0.533896
3	-0.1283	-0.169283	0.533896
4	-0.1283	-0.169283	1.489856

	AdjustmentsToCreditRating	HandsetPrice
0	-0.140707	-0.864858
1	-0.140707	-0.864858
2	-0.140707	-0.368174
3	-0.140707	-1.195980
4	2.469282	-1.195980

[5 rows x 35 columns]

We will compute the Euclidean distance between two examples in our data. In other words, our vectors  $A$  and  $B$  will be two distinct *rows* of our DataFrame `df_numerical` (which we filtered to include only numerical columns).

The code cell below randomly samples two rows from the `df_numerical` dataset and stores each in new DataFrame objects named  $A$  and  $B$ , respectively.

```
[17]: A = df_numerical.sample(replace=False)
      B = df_numerical.sample(replace=False)
```

```
[18]: A
```

```
[18]:      MonthlyRevenue  MonthlyMinutes  TotalRecurringCharge \
31056      0.450391      0.304505      0.552229
```



```

DirectorAssistedCalls  OverageMinutes  RoamingCalls  PercChangeMinutes  \
31056                0.154708          0.299959        -0.125914        -1.135682

PercChangeRevenues  DroppedCalls  BlockedCalls  ...  HandsetModels  \
31056          -0.591492          2.0222          0.019579  ...          0.487071

CurrentEquipmentDays  AgeHH1  AgeHH2  RetentionCalls  \
31056          -0.735013  0.48256  0.787922          -0.180167

RetentionOffersAccepted  ReferralsMadeBySubscriber  IncomeGroup  \
31056                -0.1283                -0.169283          0.852549

AdjustmentsToCreditRating  HandsetPrice
31056                -0.140707          -0.864858

[1 rows x 35 columns]

```

[19]:

B

[19]:

```

MonthlyRevenue  MonthlyMinutes  TotalRecurringCharge  \
42903          -0.569226          -0.491923                -0.286393

DirectorAssistedCalls  OverageMinutes  RoamingCalls  PercChangeMinutes  \
42903          -0.401714          -0.414422          0.200012        -0.626969

PercChangeRevenues  DroppedCalls  BlockedCalls  ...  HandsetModels  \
42903          0.123614          0.330446          -0.217933  ...        -0.616775

CurrentEquipmentDays  AgeHH1  AgeHH2  RetentionCalls  \
42903          -0.794114  0.754122  1.122214          -0.180167

RetentionOffersAccepted  ReferralsMadeBySubscriber  IncomeGroup  \
42903                -0.1283                -0.169283          1.489856

AdjustmentsToCreditRating  HandsetPrice
42903                -0.140707          0.790755

[1 rows x 35 columns]

```

### 1.3 Step 2: Compute the Euclidean Distance Between Two Vectors Using Python

We will first implement a function that finds the Euclidean distance in Python. Since we will be working with Python, let us convert DataFrames A and B into Python lists.

[20]:

```

list_A = A.values.flatten().tolist()
list_B = B.values.flatten().tolist()

```

```
list_A
```

```
[20]: [0.4503910267294967,
       0.3045045869325048,
       0.5522291485247021,
       0.15470821146877115,
       0.2999593842618582,
       -0.1259135548660268,
       -1.1356823309402944,
       -0.5914920040818883,
       2.022200331268707,
       0.019579057214878026,
       0.6099169936193746,
       0.4770327624039685,
       -0.2557966673525969,
       1.0943969852329347,
       -0.4850410632895274,
       -0.4307108741459736,
       0.3377959655294201,
       -0.06847081268218727,
       1.2113042357031023,
       -0.020662564533263195,
       -0.3295397355456733,
       -0.6894117553078604,
       0.38242109864799056,
       -0.5245829903418487,
       0.14600359624956538,
       0.4870710798513511,
       -0.7350126526839595,
       0.482559657577787,
       0.7879217820917835,
       -0.18016687925557376,
       -0.12830030819753901,
       -0.16928338528385997,
       0.8525494747501462,
       -0.14070742066769315,
       -0.8648576341987015]
```

Using the definition above, complete the function below that returns the Euclidean distance between its two list inputs.

You will use a traditional for loop to handle the computation for each pair of i-th coordinates of the two input lists (You can think of each pair as a 'column' in a DataFrame with just two rows -- A and B.).

Tip: to compute the square root, use the Python `math.sqrt()` function.

### 1.3.1 Graded Cell

The cell below will be graded. Remove the line "raise NotImplementedError()" before writing your code.

```

[21]: def euclidean_distance(vector1 , vector2):
    ## the sum_squares variable will contain the current value of the sum of
    →squares of each i-th coordinate pair
    sum_squares = 0

    numberOfIterations = len(vector1)

    ## TODO: Complete loop below ##

    # The number of times the loop will be executed is the length of the
    →vectors.
    #
    # At each loop iteration, you will:
    # Step 1. index into each vector and find the difference between the ith
    →element in vector2 and vector1
    # Step 2. square the difference
    # Step 3. update the value of the 'sum_squares' variable by adding the
    →result in Step 2 to
    # the existing value of sum_squares

    for i in range(numberOfIterations):

        # Inside this loop follow steps 1-3 to update the value of the
        →'sum_squares' variable by
        # adding the squared difference of the i'th coordinate pair to the sum.

        # YOUR CODE HERE
        difference=vector2[i]-vector1[i]
        square=difference**2
        sum_squares+=square

    ### TODO: Compute the Distance ###

    # Compute the square root of the variable 'sum_squares' and assign
    # that result to a new variable named 'distance'

    # YOUR CODE HERE
    distance=math.sqrt(sum_squares)

    # return the Euclidean distance
    return distance

```

### 1.3.2 Self-Check

Run the cell below to test the correctness of your code above before submitting for grading. Do not add code or delete code in the cell.

```
[22]: # Run this self-test cell to check your code;
      # do not add code or delete code in this cell
      from jn import testFunction

      try:
          p, err = testFunction(euclidean_distance)
          print(err)
      except Exception as e:
          print("Error!\n" + str(e))
```

Correct!

The code cell below tests your function. Run the cell to view the results.

```
[23]: euclidean_distance(list_A, list_B)
```

```
[23]: 4.345778982532897
```

## 1.4 Step 3: Compute the Euclidean Distance Between Two Vectors Using NumPy

The NumPy package provides an easy way to compute the Euclidean distance between two vectors. NumPy has a `norm()` function, which is part of a linear algebra module called `linalg`. You can call the function using this syntax: `np.linalg.norm()`. The `norm([vector_name])` finds a vector norm. A vector has both magnitude and direction, and calculating the vector norm finds the magnitude.

By default, the `norm()` function calculates the L2 norm, also known as the Euclidean norm since it calculates the Euclidean distance. We can therefore use the `norm()` function to calculate the distance between two vectors.

The `norm()` function requires that its input vectors be of type NumPy array. The code cell below converts DataFrame A and B to NumPy arrays and uses the `norm()` function to find the Euclidean distance.

Run the cell below and compare the results. Is the Euclidean distance the same value as what your function `euclidean_distance` produces? Try using the `norm()` function to find the Euclidean distance between the vectors `vector_A` and `vector_B` as well.

```
[24]: array1 = np.array(A)
      array2 = np.array(B)
      np.linalg.norm(array2-array1)
```

```
[24]: 4.345778982532897
```

You can see how easy it is to find the Euclidean distance between two vectors, or examples using NumPy! For more information about the `norm()` function, consult the online [documentation](#).