

Lab3

August 8, 2023

1 Lab 3: Training Decision Tree & KNN Classifiers

```
[1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
pd.options.mode.chained_assignment = None

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
```

In this Lab session, you will implement the following steps:

1. Load the Airbnb "listings" data set
2. Convert categorical features to one-hot encoded values
3. Split the data into training and test sets
4. Fit a Decision Tree classifier and evaluate the accuracy
 - Plot the accuracy of the DT model as a function of hyperparameter max depth
5. Fit a KNN classifier and evaluate the accuracy
 - Plot the accuracy of the KNN model as a function of hyperparameter k

1.1 Part 1. Load the Dataset

We will work with a preprocessed version of the Airbnb NYC "listings" data set.

Task: load the data set into a Pandas DataFrame variable named `df`:

```
[2]: # Do not remove or edit the line below:
filename = os.path.join(os.getcwd(), "data", "airbnb.csv.gz")

# YOUR CODE HERE
df=pd.read_csv(filename)
```

```
[3]: df.shape
```

```
[3]: (28022, 44)
```

```
[4]: df.head(10)
```

```
[4]:   host_response_rate  host_acceptance_rate  host_is_superhost \
0          0.800000          0.170000          False
1          0.090000          0.690000          False
2          1.000000          0.250000          False
3          1.000000          1.000000          False
4          0.890731          0.768297          False
5          1.000000          1.000000           True
6          1.000000          1.000000          False
7          1.000000          1.000000          False
8          1.000000          0.000000          False
9          1.000000          0.990000           True

      host_listings_count  host_total_listings_count  host_has_profile_pic \
0                8.0                8.0                True
1                1.0                1.0                True
2                1.0                1.0                True
3                1.0                1.0                True
4                1.0                1.0                True
5                3.0                3.0                True
6                1.0                1.0                True
7                3.0                3.0                True
8                2.0                2.0                True
9                1.0                1.0                True

      host_identity_verified  neighbourhood_group_cleansed  room_type \
0                True          Manhattan  Entire home/apt
1                True          Brooklyn  Entire home/apt
2                True          Brooklyn  Entire home/apt
3                False          Manhattan   Private room
4                True          Manhattan   Private room
5                True          Brooklyn   Private room
6                True          Brooklyn  Entire home/apt
7                True          Manhattan   Private room
8                True          Brooklyn   Private room
9                True          Brooklyn  Entire home/apt

      accommodates  ...  review_scores_communication  review_scores_location \
0                1  ...                4.79                4.86
1                3  ...                4.80                4.71
2                4  ...                5.00                4.50
3                2  ...                4.42                4.87
4                1  ...                4.95                4.94
```

5	2	...	4.82	4.87
6	3	...	4.80	4.67
7	1	...	4.95	4.84
8	1	...	5.00	5.00
9	4	...	4.91	4.93

	review_scores_value	instant_bookable	calculated_host_listings_count	\
0	4.41	False		3
1	4.64	False		1
2	5.00	False		1
3	4.36	False		1
4	4.92	False		1
5	4.73	False		3
6	4.57	True		1
7	4.84	True		1
8	5.00	False		2
9	4.78	True		2

	calculated_host_listings_count_entire_homes	\
0	3	
1	1	
2	1	
3	0	
4	0	
5	1	
6	1	
7	0	
8	0	
9	1	

	calculated_host_listings_count_private_rooms	\
0	0	
1	0	
2	0	
3	1	
4	1	
5	2	
6	0	
7	1	
8	2	
9	1	

	calculated_host_listings_count_shared_rooms	reviews_per_month	\
0	0	0.33	
1	0	4.86	
2	0	0.02	
3	0	3.68	

4	0	0.87
5	0	1.48
6	0	1.24
7	0	1.82
8	0	0.07
9	0	3.05

	n_host_verifications
0	9
1	6
2	3
3	4
4	7
5	7
6	7
7	5
8	5
9	8

[10 rows x 44 columns]

```
[5]: df.columns
```

```
[5]: Index(['host_response_rate', 'host_acceptance_rate', 'host_is_superhost',
        'host_listings_count', 'host_total_listings_count',
        'host_has_profile_pic', 'host_identity_verified',
        'neighbourhood_group_cleansed', 'room_type', 'accommodates',
        'bathrooms', 'bedrooms', 'beds', 'amenities', 'price', 'minimum_nights',
        'maximum_nights', 'minimum_minimum_nights', 'maximum_minimum_nights',
        'minimum_maximum_nights', 'maximum_maximum_nights',
        'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', 'has_availability',
        'availability_30', 'availability_60', 'availability_90',
        'availability_365', 'number_of_reviews', 'number_of_reviews_ltm',
        'number_of_reviews_l30d', 'review_scores_rating',
        'review_scores_cleanliness', 'review_scores_checkin',
        'review_scores_communication', 'review_scores_location',
        'review_scores_value', 'instant_bookable',
        'calculated_host_listings_count',
        'calculated_host_listings_count_entire_homes',
        'calculated_host_listings_count_private_rooms',
        'calculated_host_listings_count_shared_rooms', 'reviews_per_month',
        'n_host_verifications'],
        dtype='object')
```

1.2 Part 2. One-Hot Encode Categorical Values

Transform the string-valued categorical features into numerical boolean values using one-hot encoding.

1.2.1 a. Find the Columns Containing String Values

First, let us identify all features that need to be one-hot encoded:

```
[6]: df.dtypes
```

```
[6]: host_response_rate      float64
     host_acceptance_rate    float64
     host_is_superhost        bool
     host_listings_count      float64
     host_total_listings_count float64
     host_has_profile_pic      bool
     host_identity_verified    bool
     neighbourhood_group_cleansed object
     room_type                 object
     accommodates              int64
     bathrooms                 float64
     bedrooms                  float64
     beds                     float64
     amenities                 object
     price                     float64
     minimum_nights            int64
     maximum_nights            int64
     minimum_minimum_nights    float64
     maximum_minimum_nights    float64
     minimum_maximum_nights    float64
     maximum_maximum_nights    float64
     minimum_nights_avg_ntm    float64
     maximum_nights_avg_ntm    float64
     has_availability           bool
     availability_30            int64
     availability_60            int64
     availability_90            int64
     availability_365           int64
     number_of_reviews          int64
     number_of_reviews_ltm      int64
     number_of_reviews_l30d     int64
     review_scores_rating       float64
     review_scores_cleanliness  float64
     review_scores_checkin      float64
     review_scores_communication float64
     review_scores_location     float64
     review_scores_value        float64
     instant_bookable           bool
     calculated_host_listings_count int64
     calculated_host_listings_count_entire_homes int64
     calculated_host_listings_count_private_rooms int64
     calculated_host_listings_count_shared_rooms int64
     reviews_per_month          float64
```

```
n_host_verifications          int64
dtype: object
```

Task: add all of the column names of variables of type 'object' to a list named `to_encode`

```
[7]: # YOUR CODE HERE
to_encode=list(df.select_dtypes(include=['object']).columns)
```

Let's take a closer look at the candidates for one-hot encoding

```
[8]: df[to_encode].nunique()
```

```
[8]: neighbourhood_group_cleansed    5
room_type                          4
amenities                        25020
dtype: int64
```

Notice that one column stands out as containing too many values for us to attempt to transform. For this exercise, the best choice is to simply remove this column. Of course, this means losing potentially useful information. In a real-life situation, you would want to retain all of the information in a column, or you could selectively keep information in.

In the code cell below, drop this column from Dataframe `df` and from the `to_encode` list.

```
[9]: # YOUR SOLUTION HERE
column_to_drop='amenities'
df=df.drop(column_to_drop,axis=1)
to_encode.remove(column_to_drop)
```

1.2.2 b. One-Hot Encode all Unique Values

All of the other columns in `to_encode` have reasonably small numbers of unique values, so we are going to simply one-hot encode every unique value of those columns.

Task: complete the code below to create one-hot encoded columns **Tip:** Use the `sklearn` `OneHotEncoder` class

```
[10]: from sklearn.preprocessing import OneHotEncoder

# Create the encoder:
encoder = OneHotEncoder(sparse=False)

# Apply the encoder:
df_enc = pd.DataFrame(encoder.fit_transform(df[to_encode]))

# Reinstate the original column names:
df_enc.columns = encoder.get_feature_names(to_encode)
```

```
[11]: df_enc.head()
```

```
[11]: neighbourhood_group_cleansed_Bronx  neighbourhood_group_cleansed_Brooklyn  \
0                                     0.0                                     0.0
1                                     0.0                                     1.0
2                                     0.0                                     1.0
3                                     0.0                                     0.0
```

4	0.0	0.0
---	-----	-----

	neighbourhood_group_cleansed_Manhattan \
0	1.0
1	0.0
2	0.0
3	1.0
4	1.0

	neighbourhood_group_cleansed_Queens \
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

	neighbourhood_group_cleansed_Staten Island	room_type_Entire home/apt \
0	0.0	1.0
1	0.0	1.0
2	0.0	1.0
3	0.0	0.0
4	0.0	0.0

	room_type_Hotel room	room_type_Private room	room_type_Shared room
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	1.0	0.0
4	0.0	1.0	0.0

Task: You can now remove the original columns that we have just transformed from DataFrame df.

```
[12]: # YOUR CODE HERE
df.drop(columns=to_encode,inplace=True)
```

```
[13]: df.head()
```

	host_response_rate	host_acceptance_rate	host_is_superhost \
0	0.800000	0.170000	False
1	0.090000	0.690000	False
2	1.000000	0.250000	False
3	1.000000	1.000000	False
4	0.890731	0.768297	False

	host_listings_count	host_total_listings_count	host_has_profile_pic \
0	8.0	8.0	True
1	1.0	1.0	True
2	1.0	1.0	True

3	1.0	1.0	True
4	1.0	1.0	True

	host_identity_verified	accommodates	bathrooms	bedrooms	...	\
0	True	1	1.0	1.323567	...	
1	True	3	1.0	1.000000	...	
2	True	4	1.5	2.000000	...	
3	False	2	1.0	1.000000	...	
4	True	1	1.0	1.000000	...	

	review_scores_communication	review_scores_location	review_scores_value	\
0	4.79	4.86	4.41	
1	4.80	4.71	4.64	
2	5.00	4.50	5.00	
3	4.42	4.87	4.36	
4	4.95	4.94	4.92	

	instant_bookable	calculated_host_listings_count	\
0	False	3	
1	False	1	
2	False	1	
3	False	1	
4	False	1	

	calculated_host_listings_count_entire_homes	\
0	3	
1	1	
2	1	
3	0	
4	0	

	calculated_host_listings_count_private_rooms	\
0	0	
1	0	
2	0	
3	1	
4	1	

	calculated_host_listings_count_shared_rooms	reviews_per_month	\
0	0	0.33	
1	0	4.86	
2	0	0.02	
3	0	3.68	
4	0	0.87	

	n_host_verifications
0	9

1	6
2	3
3	4
4	7

[5 rows x 41 columns]

Task: You can now join the transformed categorical features contained in `df_enc` with DataFrame `df`

```
[14]: # YOUR CODE HERE
df=pd.concat([df,df_enc],axis=1)
```

Glance at the resulting column names:

```
[15]: df.columns
```

```
[15]: Index(['host_response_rate', 'host_acceptance_rate', 'host_is_superhost',
        'host_listings_count', 'host_total_listings_count',
        'host_has_profile_pic', 'host_identity_verified', 'accommodates',
        'bathrooms', 'bedrooms', 'beds', 'price', 'minimum_nights',
        'maximum_nights', 'minimum_minimum_nights', 'maximum_minimum_nights',
        'minimum_maximum_nights', 'maximum_maximum_nights',
        'minimum_nights_avg_ntm', 'maximum_nights_avg_ntm', 'has_availability',
        'availability_30', 'availability_60', 'availability_90',
        'availability_365', 'number_of_reviews', 'number_of_reviews_ltm',
        'number_of_reviews_l30d', 'review_scores_rating',
        'review_scores_cleanliness', 'review_scores_checkin',
        'review_scores_communication', 'review_scores_location',
        'review_scores_value', 'instant_bookable',
        'calculated_host_listings_count',
        'calculated_host_listings_count_entire_homes',
        'calculated_host_listings_count_private_rooms',
        'calculated_host_listings_count_shared_rooms', 'reviews_per_month',
        'n_host_verifications', 'neighbourhood_group_cleansed_Bronx',
        'neighbourhood_group_cleansed_Brooklyn',
        'neighbourhood_group_cleansed_Manhattan',
        'neighbourhood_group_cleansed_Queens',
        'neighbourhood_group_cleansed_Staten Island',
        'room_type_Entire home/apt', 'room_type_Hotel room',
        'room_type_Private room', 'room_type_Shared room'],
        dtype='object')
```

Check for missing values.

```
[16]: # YOUR CODE HERE
missing_values=df.isnull().sum()
```

1.3 Part 3. Create Training and Test Data Sets

1.3.1 a. Create Labeled Examples

Task: Choose columns from our data set to create labeled examples.

In the airbnb dataset, we will choose column `host_is_superhost` to be the label. The remaining columns will be the features.

Obtain the features from DataFrame `df` and assign to `X`. Obtain the label from DataFrame `df` and assign to `Y`

```
[17]: # YOUR CODE HERE
X=df.drop('host_is_superhost',axis=1)
y=df['host_is_superhost']

[18]: print("Number of examples: " + str(X.shape[0]))
print("\nNumber of Features:" + str(X.shape[1]))
print(str(list(X.columns)))
```

Number of examples: 28022

Number of Features:49

```
['host_response_rate', 'host_acceptance_rate', 'host_listings_count',
'host_total_listings_count', 'host_has_profile_pic', 'host_identity_verified',
'accommodates', 'bathrooms', 'bedrooms', 'beds', 'price', 'minimum_nights',
'maximum_nights', 'minimum_minimum_nights', 'maximum_minimum_nights',
'minimum_maximum_nights', 'maximum_maximum_nights', 'minimum_nights_avg_ntm',
'maximum_nights_avg_ntm', 'has_availability', 'availability_30',
'availability_60', 'availability_90', 'availability_365', 'number_of_reviews',
'number_of_reviews_ltm', 'number_of_reviews_l30d', 'review_scores_rating',
'review_scores_cleanliness', 'review_scores_checkin',
'review_scores_communication', 'review_scores_location', 'review_scores_value',
'instant_bookable', 'calculated_host_listings_count',
'calculated_host_listings_count_entire_homes',
'calculated_host_listings_count_private_rooms',
'calculated_host_listings_count_shared_rooms', 'reviews_per_month',
'n_host_verifications', 'neighbourhood_group_cleansed_Bronx',
'neighbourhood_group_cleansed_Brooklyn',
'neighbourhood_group_cleansed_Manhattan', 'neighbourhood_group_cleansed_Queens',
'neighbourhood_group_cleansed_Staten Island', 'room_type_Entire home/apt',
'room_type_Hotel room', 'room_type_Private room', 'room_type_Shared room']
```

1.3.2 b. Split Examples into Training and Test Sets

Task: In the code cell below create training and test sets out of the labeled examples using Scikit-learn's `train_test_split()` function.

Specify: * A test set that is one third (.33) of the size of the data set. * A seed value of '123'.

```
[26]: # YOUR CODE HERE
test_size=0.33
random_seed=123
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=test_size,random_state=random_see
```

Check that the dimensions of the training and test datasets are what you expected

```
[27]: print(X_train.shape)
      print(X_test.shape)
```

```
(18774, 49)
```

```
(9248, 49)
```

1.4 Part 4. Implement a Decision Tree Classifier

The code cell below contains a shell of a function named `train_test_DT()`. This function should train a Decision Tree classifier on the training data, test the resulting model on the test data, and compute and return the accuracy score of the resulting predicted class labels on the test data. Remember to use `DecisionTreeClassifier()` to create a model object.

Task: Complete the function to make it work.

```
[32]: def train_test_DT(X_train, X_test, y_train, y_test, leaf, depth,
    ↪crit='entropy'):
    """
    Fit a Decision Tree classifier to the training data X_train, y_train.
    Return the accuracy of resulting predictions on the test set.
    Parameters:
        leaf := The minimum number of samples required to be at a leaf node
        depth := The maximum depth of the tree
        crit := The function to be used to measure the quality of a split.
    ↪Default: gini.
    """

    # YOUR CODE HERE

    ↪
    ↪model=DecisionTreeClassifier(min_samples_leaf=leaf,max_depth=depth,criterion=crit)
    model.fit(X_train,y_train)
    y_pred=model.predict(X_test)
    acc_score=model.score(X_test,y_test)

    return acc_score
```

Visualization The cell below contains a function that you will use to compare the accuracy results of training multiple models with different hyperparameter values.

Function `visualize_accuracy()` accepts two arguments: 1. a list of hyperparameter values 2. a list of accuracy scores

Both lists must be of the same size.

```
[33]: # Do not remove or edit the code below

def visualize_accuracy(hyperparam_range, acc):
```

```

fig = plt.figure()
ax = fig.add_subplot(111)
p = sns.lineplot(x=hyperparam_range, y=acc, marker='o', label = 'Full_
→training set')

plt.title('Test set accuracy of the model predictions, for ' + ', '.
→join([str(h) for h in hyperparam_range]))
ax.set_xlabel('Hyperparameter value')
ax.set_ylabel('Accuracy')
plt.show()

```

Train on Different Values of Hyperparameter Max Depth Task:

Complete function `train_multiple_trees()` in the code cell below. The function should train multiple decision trees and return a list of accuracy scores.

The function will:

1. accept list `max_depth_range` and `leaf` as parameters; list `max_depth_range` will contain multiple values for hyperparameter max depth.
2. loop over list `max_depth_range` and at each iteration:
 - a. index into list `max_depth_range` to obtain a value for max depth
 - b. call `train_test_DT` with the training and test set, the value of max depth, and the value of leaf
 - c. print the resulting accuracy score
 - d. append the accuracy score to list `accuracy_list`

```

[34]: def train_multiple_trees(max_depth_range, leaf):

    accuracy_list = []

    # YOUR CODE HERE
    for max_depth in max_depth_range:
        acc_score=train_test_DT(X_train,X_test,y_train,y_test,leaf,max_depth)
        print(f"Max Depth: {max_depth},Accuracy:{acc_score}")
        accuracy_list.append(acc_score)
    return accuracy_list

```

The code cell below tests function `train_multiple_trees()` and calls function `visualize_accuracy()` to visualize the results.

```


[35]: max_depth_range = [8, 32]
leaf = 1

acc = train_multiple_trees(max_depth_range, leaf)

visualize_accuracy(max_depth_range, acc)

```

Max Depth: 8,Accuracy:0.8333693771626297
Max Depth: 32,Accuracy:0.8004974048442907



Analysis: Is this graph conclusive for determining a good value of max depth?

No, this graph does not have enough data points to accurately determine the best value. There are only two values of max depth found-- 8 and 32. Between them, however, it is seen that a max depth of 8 yielded a higher accuracy than the max depth of 32.


Task: Let's train on more values for max depth.

In the code cell below:

1. call `train_multiple_trees()` with arguments `max_depth_range` and `leaf`
2. call `visualize_accuracy()` with arguments `max_depth_range` and `acc`

```
[48]: max_depth_range = [2**i for i in range(6)]  
leaf = 1  
acc = train_multiple_trees(max_depth_range,leaf)  
  
visualize_accuracy(max_depth_range,acc)
```

Max Depth: 1,Accuracy:0.7563797577854672
Max Depth: 2,Accuracy:0.7563797577854672
Max Depth: 4,Accuracy:0.810878027681661
Max Depth: 8,Accuracy:0.8331531141868512
Max Depth: 16,Accuracy:0.811310553633218
Max Depth: 32,Accuracy:0.8066608996539792



Analysis: Analyze this graph. Keep in mind that this is the performance on the test set, and pay attention to the scale of the y-axis. Answer the following questions in the cell below. How would you go about choosing the best model based on this plot? Is it conclusive? What other hyperparameters of interest would you want to vary to make sure you are finding the best model fit?

This graph shows that the best model performance is at a max depth of 8 with the highest accuracy score of 0.807, but it is not conclusive. To figure out the best model more conclusively, other hyperparameters such as `min_samples_leaf`, `max_leaf_nodes`, or `crit.` would be useful to vary. Varying these hyperparameters and then seeing their effects on model performance would let you see the combination that maximizes accuracy on unseen data.

1.5 Part 5. Implement a KNN Classifier

Note: In this section you will train KNN classifiers using the same training and test data.

The code cell below contains a shell of a function named `train_test_knn()`. This function should train a KNN classifier on the training data, test the resulting model on the test data, and compute and return the accuracy score of the resulting predicted class labels on the test data.

Remember to use `KNeighborsClassifier()` to create a model object and call the method with one parameter: `n_neighbors = k`.

Task: Complete the function to make it work.

```
[49]: def train_test_knn(X_train, X_test, y_train, y_test, k):
      '''
      Fit a k Nearest Neighbors classifier to the training data X_train, y_train.
      Return the accuracy of resulting predictions on the test data.
```

```

'''

# YOUR CODE HERE
model=KNeighborsClassifier(n_neighbors=k)
model.fit(X_train,y_train)
acc_score=model.score(X_test,y_test)

return acc_score

```

Train on Different Values of Hyperparameter K Task:

Just as you did above, complete function `train_multiple_knns()` in the code cell below. The function should train multiple KNN models and return a list of accuracy scores.

The function will:

1. accept list `k_range` as a parameter; this list will contain multiple values for hyperparameter `k`
2. loop over list `k_range` and at each iteration:
 - a. index into list `k_range` to obtain a value for `k`
 - b. call `train_test_knn` with the training and test set, and the value of `k`
 - c. print the resulting accuracy score
 - d. append the accuracy score to list `accuracy_list`

```

[50]: def train_multiple_knns(k_range):

    accuracy_list = []

    # YOUR CODE HERE
    for k in k_range:
        acc_score=train_test_knn(X_train,X_test,y_train,y_test,k)
        print(f"K:{k},Accuracy:{acc_score}")
        accuracy_list.append(acc_score)

    return accuracy_list

```

The code cell below uses your `train_multiple_knn()` function to train 3 KNN models, specifying three values for `k`: 3, 30, and 300. It calls function `visualize_accuracy()` to visualize the results. Note: this make take a second.

```

[51]: k_range = [3, 30, 300]
    acc = train_multiple_knns(k_range)


    visualize_accuracy(k_range, acc)

```

```

K:3,Accuracy:0.759515570934256
K:30,Accuracy:0.7755190311418685
K:300,Accuracy:0.7696799307958477

```



Task: Let's train on more values for k
In the code cell below:

1. call `train_multiple_knns()` with argument `k_range`
2. call `visualize_accuracy()` with arguments `k_range` and the resulting accuracy list obtained from `train_multiple_knns()`

```
[54]: k_range = np.arange(1, 40, step = 3)

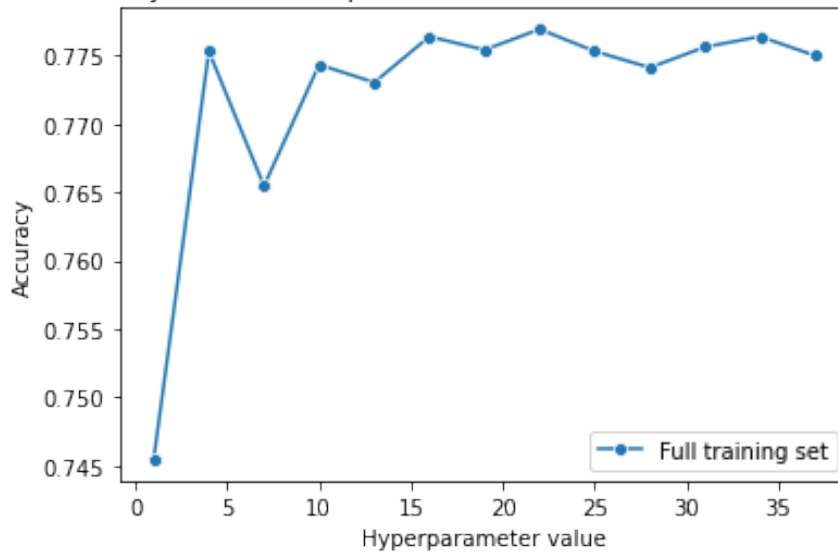
# YOUR CODE HERE
acc_knn=train_multiple_knns(k_range)
visualize_accuracy(k_range,acc_knn)
```

```
K:1,Accuracy:0.7454584775086506
K:4,Accuracy:0.77530276816609
K:7,Accuracy:0.7654628027681661
K:10,Accuracy:0.7743295847750865
K:13,Accuracy:0.7730320069204152
K:16,Accuracy:0.7763840830449827
K:19,Accuracy:0.7754108996539792
K:22,Accuracy:0.776924740484429
K:25,Accuracy:0.77530276816609
K:28,Accuracy:0.7741133217993079
K:31,Accuracy:0.7756271626297578
```


K:34,Accuracy:0.7763840830449827

K:37,Accuracy:0.7749783737024222

Test set accuracy of the model predictions, for 1,4,7,10,13,16,19,22,25,28,31,34,37



Analysis: Compare the performance of the KNN model relative to the Decision Tree model, with various hyperparameter values and record your findings in the cell below.

The KNN model shows generally stable performance accuracy across K values, with only a small variation in accuracy scores. The decision Tree model shows accuracy scores ranging from 0.756 to 0.833 in comparison to KNNs 0.745 to 0.777. Because of this, the Decision Tree model outperformed the KNN model when considering the accuracy on this dataset.

[]:	
[]:	
[]:	
[]:	
[]:	