

# MissingData

August 8, 2023

## 1 Finding and Replacing Missing Data

```
[1]: import pandas as pd
import numpy as np
import os
```

### 1.1 Load and Inspect the Data

```
[2]: filename = os.path.join(os.getcwd(), "data", "adult.data.partial.missing")
df = pd.read_csv(filename, header=0)
```

```
[3]: df.shape
```

```
[3]: (7000, 15)
```

```
[4]: df.head()
```

```
[4]:   age  workclass  fnlwgt  education  education-num  marital-status \
0  36.0   State-gov  112074   Doctorate           16   Never-married
1  35.0    Private   32528    HS-grad            9  Married-civ-spouse
2  21.0    Private  270043  Some-college          10   Never-married
3  45.0    Private  168837  Some-college          10  Married-civ-spouse
4  39.0    Private  297449   Bachelors           13  Married-civ-spouse
```

```
   occupation  relationship  race  sex_selfID  capital-gain \
0  Prof-specialty  Not-in-family  White  Non-Female           0
1  Handlers-cleaners      Husband  White  Non-Female           0
2   Other-service    Own-child  White    Female           0
3   Adm-clerical      Wife  White    Female           0
4  Prof-specialty      Husband  White  Non-Female           0
```

```
   capital-loss  hours-per-week  native-country  label
0              0             45.0  United-States  <=50K
1              0             45.0  United-States  <=50K
2              0             16.0  United-States  <=50K
3              0             24.0      Canada    >50K
4              0             40.0  United-States    >50K
```

## 1.2 Dealing with Missing Data

Our goal will be to identify which columns in a dataset have missing values, and to replace a missing value in a column with the mean of the other values in that column. We will add dummy variables to our dataset to indicate which columns initially had missing values.

### 1.2.1 Step 1: Identify Missing Values Using Pandas `isnull()` Method

First let us check if there are missing values in DataFrame `df`.

```
[5]: df.isnull().values.any()
```

```
[5]: True
```

DataFrame `df` contains missing values! The Pandas `isnull()` method returns `True/False` values indicating whether a value is or is not missing in a particular position in a DataFrame or Series. This method recognizes various spellings of missingness like `NaN`, `nan`, `None`, and `NA` among others. Consult the online [documentation](#) for more information.

```
[6]: df.isnull().head()
```

```
[6]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

  

	occupation	relationship	race	sex_selfID	capital-gain	capital-loss	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

  

	hours-per-week	native-country	label
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

The code cell below counts the number of times a missing value occurs in each column. It applies the `isnull()` method and then aggregates the results by columns using the `np.sum()` function. For more information about `np.sum()`, consult the online [documentation](#).

```
[7]: nan_count = np.sum(df.isnull(), axis = 0)
nan_count
```

```
[7]: age          35
workclass      375
fnlwgt         0
education      0
```

```

education-num      0
marital-status     0
occupation         375
relationship       0
race               0
sex_selfID        0
capital-gain       0
capital-loss       0
hours-per-week     70
native-country     138
label             0
dtype: int64

```

The code cell below stores the names of the columns with detected missing values into a Python list.

```

[8]: condition = nan_count != 0 # look for all columns with missing values

col_names = nan_count[condition].index # get the column names
print(col_names)

nan_cols = list(col_names) # convert column names to list
print(nan_cols)

```

```

Index(['age', 'workclass', 'occupation', 'hours-per-week', 'native-country'],
dtype='object')
['age', 'workclass', 'occupation', 'hours-per-week', 'native-country']

```

### 1.2.2 Step 2: Choose Which Values to Fill

We can see that five columns in our DataFrame contain missing values. Would you want to replace the missing values with something for every one of these columns?

Let's take a look at the data types of the columns that contain missing values using dtypes.

```

[9]: nan_col_types = df[nan_cols].dtypes
nan_col_types

```

```

[9]: age          float64
workclass        object
occupation        object
hours-per-week   float64
native-country    object
dtype: object

```

For three of the five identified columns, the type is 'object'. Is this a problem? A common approach to dealing with the missing values is to replace those values with either the mean, the median, or some other type of 'representative' value wherever a nan occurs. This, of course, assumes that the column is numerical to begin with. That does not seem to be true for the workclass, occupation, and native-country variables. Let us confirm:

```
[10]: print(df['workclass'].unique())
print(df['occupation'].unique())
print(df['native-country'].unique())
```

```
['State-gov' 'Private' nan 'Self-emp-not-inc' 'Local-gov' 'Self-emp-inc'
 'Federal-gov' 'Without-pay']
['Prof-specialty' 'Handlers-cleaners' 'Other-service' 'Adm-clerical'
 'Craft-repair' 'Sales' nan 'Exec-managerial' 'Farming-fishing'
 'Machine-op-inspct' 'Transport-moving' 'Tech-support' 'Priv-house-serv'
 'Protective-serv' 'Armed-Forces']
['United-States' 'Canada' 'England' 'Germany' 'Cuba' nan 'Puerto-Rico'
 'Mexico' 'Nicaragua' 'China' 'South' 'India' 'Vietnam' 'Philippines'
 'El-Salvador' 'Guatemala' 'Japan' 'Jamaica' 'Peru' 'France' 'Greece'
 'Italy' 'Columbia' 'Honduras' 'Iran' 'Poland' 'Haiti'
 'Dominican-Republic' 'Scotland' 'Yugoslavia' 'Trinidad&Tobago' 'Ireland'
 'Portugal' 'Taiwan' 'Hong' 'Ecuador' 'Laos' 'Hungary' 'Thailand'
 'Outlying-US(Guam-USVI-etc)' 'Cambodia']
```

The concept of ‘mean’ is not defined for string entries, so filling in the missing values with the mean of the column wouldn’t work here. In real business settings, one way to go about filling in the missing values would be to fit a model that predicts the country based on other values. All data-filling methods come with caveats, and some may threaten the validity of your larger analytical conclusions.

For the rest of this exercise, we will focus only on the numerical variables, for which it makes sense to replace every missing value with the mean of the column. Those are age and hours-per-week columns.

### 1.2.3 Step 3: Create ‘Dummy’ Variables for Missing Values

No method of imputing missing values is perfect, and for this reason it makes sense to keep track of which values we artificially created.

The code cell below looks at the the values in columns age and hours-per-week and stores the corresponding True/False values (True if the value is missing and False if the value is present) in new columns age\_na and hours-per-week\_na. Run the cell and inspect the new columns.

```
[11]: df['age_na'] = df['age'].isnull()
df['hours-per-week_na'] = df['hours-per-week'].isnull()
df.head()
```

```
[11]:   age  workclass  fnlwgt   education  education-num   marital-status \
0  36.0   State-gov  112074   Doctorate             16   Never-married
1  35.0    Private   32528    HS-grad              9  Married-civ-spouse
2  21.0    Private  270043  Some-college             10   Never-married
3  45.0    Private  168837  Some-college             10  Married-civ-spouse
4  39.0    Private  297449   Bachelors              13  Married-civ-spouse

      occupation  relationship   race  sex_selfID  capital-gain \
0  Prof-specialty  Not-in-family  White  Non-Female           0
1  Handlers-cleaners      Husband  White  Non-Female           0
```

2	Other-service	Own-child	White	Female	0
3	Adm-clerical	Wife	White	Female	0
4	Prof-specialty	Husband	White	Non-Female	0

  

	capital-loss	hours-per-week	native-country	label	age_na \
0	0	45.0	United-States	<=50K	False
1	0	45.0	United-States	<=50K	False
2	0	16.0	United-States	<=50K	False
3	0	24.0	Canada	>50K	False
4	0	40.0	United-States	>50K	False

  

	hours-per-week_na
0	False
1	False
2	False
3	False
4	False

#### 1.2.4 Step 4: Fill the Missing Values Using Pandas fillna() Method

The Pandas `fillna()` method is used to "fill in" missing values in a Series or DataFrame object. Consult the online [documentation](#) for more information about how to use the `fillna()` method. The code cell below uses `fillna()` to fill in values for the missing values in the `age` column. It fills in the missing values with the mean value of all of the existing values in the that column. It uses the Pandas `mean()` method to compute the replacement values. For more information about `mean()`, consult the online [documentation](#).

Tip: when working with `fillna()`, make sure that you do not just create a copy object with the filled values, but change the original values of the `df` object by specifying the `inplace = True` parameter value.

First inspect some of the columns that contain missing values.

```
[12]: df.loc[df['age'].isnull()]
```

```
[12]:   age      workclass  fnlwgt  education  education-num \
453  NaN      Private  117166    Bachelors             13
654  NaN      Private  65545     Masters             14
865  NaN  Self-emp-not-inc  93806  Some-college             10
1206 NaN      Private  441637     HS-grad              9
1262 NaN      Private  350440     HS-grad              9
1302 NaN      Private  317443  Some-college             10
1496 NaN      Private  99185     HS-grad              9
2100 NaN      Private  179271  Some-college             10
2581 NaN      Private  145160  Some-college             10
2651 NaN      Private  151580  Some-college             10
2961 NaN      Private  363219  Some-college             10
3174 NaN  Self-emp-not-inc  96245     HS-grad              9
3370 NaN      Private  214502         9th              5
3594 NaN      Private  265807  Some-college             10
```

3721	NaN	Private	322391	11th	7
3769	NaN	Local-gov	82393	HS-grad	9
3993	NaN	Private	232024	11th	7
3997	NaN	Private	235894	11th	7
4048	NaN	Private	202498	11th	7
4100	NaN	Private	33644	HS-grad	9
4253	NaN	Private	148524	HS-grad	9
4670	NaN	Self-emp-not-inc	29054	HS-grad	9
4802	NaN	Private	173208	Masters	14
4828	NaN	Private	191982	Assoc-voc	11
4866	NaN	Local-gov	286342	Masters	14
5299	NaN	Private	329426	Masters	14
5420	NaN	Federal-gov	239074	Assoc-acdm	12
5871	NaN	Private	298635	Masters	14
5949	NaN	Private	157894	Some-college	10
6007	NaN	Private	266635	HS-grad	9
6153	NaN	Private	236818	Assoc-voc	11
6219	NaN	Private	57916	HS-grad	9
6466	NaN	Private	223515	HS-grad	9
6596	NaN	Private	152307	HS-grad	9
6833	NaN	Private	289458	Bachelors	13

	marital-status	occupation	relationship \
453	Never-married	Exec-managerial	Not-in-family
654	Divorced	NaN	Own-child
865	Married-civ-spouse	Sales	Husband
1206	Married-civ-spouse	Tech-support	Husband
1262	Married-civ-spouse	Exec-managerial	Husband
1302	Never-married	Adm-clerical	Own-child
1496	Married-civ-spouse	Sales	Husband
2100	Married-civ-spouse	Craft-repair	Husband
2581	Married-civ-spouse	Machine-op-inspct	Husband
2651	Married-civ-spouse	Prof-specialty	Husband
2961	Never-married	Other-service	Not-in-family
3174	Married-civ-spouse	Machine-op-inspct	Husband
3370	Married-civ-spouse	Handlers-cleaners	Husband
3594	Separated	Craft-repair	Not-in-family
3721	Separated	Other-service	Unmarried
3769	Never-married	Handlers-cleaners	Unmarried
3993	Never-married	Machine-op-inspct	Own-child
3997	Married-civ-spouse	Exec-managerial	Husband
4048	Married-civ-spouse	Handlers-cleaners	Husband
4100	Never-married	Adm-clerical	Own-child
4253	Married-civ-spouse	Transport-moving	Husband
4670	Married-civ-spouse	Farming-fishing	Husband
4802	Married-civ-spouse	Prof-specialty	Husband
4828	Never-married	Adm-clerical	Own-child

4866	Never-married	Prof-specialty	Not-in-family
5299	Never-married	Exec-managerial	Not-in-family
5420	Married-civ-spouse	Other-service	Husband
5871	Married-civ-spouse	Prof-specialty	Husband
5949	Never-married	Other-service	Own-child
6007	Never-married	Other-service	Own-child
6153	Never-married	Prof-specialty	Unmarried
6219	Separated	Farming-fishing	Own-child
6466	Married-civ-spouse	Craft-repair	Husband
6596	Married-civ-spouse	Machine-op-inspct	Husband
6833	Never-married	Exec-managerial	Not-in-family

	race	sex_selfID	capital-gain	capital-loss	\
453	White	Non-Female	0	0	
654	White	Female	0	0	
865	White	Non-Female	0	0	
1206	White	Non-Female	0	0	
1262	Asian-Pac-Islander	Non-Female	0	0	
1302	Black	Female	0	0	
1496	White	Non-Female	7298	0	
2100	White	Non-Female	0	0	
2581	White	Non-Female	0	0	
2651	White	Non-Female	15024	0	
2961	White	Female	0	0	
3174	White	Non-Female	0	0	
3370	White	Non-Female	0	0	
3594	White	Non-Female	0	0	
3721	Black	Female	0	0	
3769	Asian-Pac-Islander	Non-Female	0	1590	
3993	White	Non-Female	0	0	
3997	White	Non-Female	0	0	
4048	White	Non-Female	0	0	
4100	White	Female	0	0	
4253	White	Non-Female	0	2057	
4670	White	Non-Female	0	0	
4802	White	Non-Female	0	0	
4828	White	Female	0	0	
4866	White	Female	0	0	
5299	White	Non-Female	0	0	
5420	White	Non-Female	0	0	
5871	Asian-Pac-Islander	Non-Female	0	0	
5949	Black	Non-Female	0	0	
6007	Black	Non-Female	0	0	
6153	Black	Female	0	0	
6219	White	Non-Female	0	0	
6466	White	Non-Female	0	0	
6596	White	Non-Female	0	0	

6833		White	Female	0	0
------	--	-------	--------	---	---

  

	hours-per-week	native-country	label	age_na	hours-per-week_na
453	50.0	United-States	<=50K	True	False
654	55.0	United-States	<=50K	True	False
865	55.0	United-States	<=50K	True	False
1206	40.0	United-States	<=50K	True	False
1262	40.0	United-States	>50K	True	False
1302	15.0	United-States	<=50K	True	False
1496	50.0	United-States	>50K	True	False
2100	50.0	United-States	>50K	True	False
2581	43.0	United-States	<=50K	True	False
2651	40.0	United-States	>50K	True	False
2961	20.0	United-States	<=50K	True	False
3174	40.0	United-States	<=50K	True	False
3370	50.0	United-States	>50K	True	False
3594	45.0	United-States	<=50K	True	False
3721	NaN	United-States	<=50K	True	True
3769	45.0	United-States	<=50K	True	False
3993	55.0	United-States	<=50K	True	False
3997	55.0	United-States	<=50K	True	False
4048	40.0	Columbia	<=50K	True	False
4100	30.0	United-States	<=50K	True	False
4253	40.0	United-States	<=50K	True	False
4670	50.0	United-States	<=50K	True	False
4802	25.0	United-States	<=50K	True	False
4828	55.0	United-States	<=50K	True	False
4866	32.0	United-States	>50K	True	False
5299	37.0	United-States	<=50K	True	False
5420	40.0	United-States	<=50K	True	False
5871	40.0	Hong	>50K	True	False
5949	20.0	United-States	<=50K	True	False
6007	30.0	United-States	<=50K	True	False
6153	26.0	United-States	<=50K	True	False
6219	40.0	United-States	<=50K	True	False
6466	45.0	United-States	<=50K	True	False
6596	40.0	United-States	<=50K	True	False
6833	40.0	United-States	<=50K	True	False

```
[13]: # look at one row that contains a missing value for age
print("Row 654: " + str(df['age'][654]))

# compute mean for all non null age values
mean_ages=df['age'].mean()
print("mean value for all age columns: " + str(mean_ages))

# fill all missing values with the mean
```



```
df['age'].fillna(value=mean_ages, inplace=True)

# look at one of the rows that contained a missing value for age.
# It should now contain the mean
print("Row 654: " + str(df['age'][654]))
```

```
Row 654: nan
mean value for all age columns: 38.61981335247667
Row 654: 38.61981335247667
```

In the code cell below, do the same for the hours-per-week column.

1. Compute the mean value of the hours-per-week column and save the result to variable `mean_hours`
2. Use `fillna` to change the values of the missing columns to `mean_hours`.

### 1.2.5 Graded Cell

The cell below will be graded. Remove the line `"raise NotImplementedError()"` before writing your code.

```
[14]: # YOUR CODE HERE
mean_hours=df['hours-per-week'].mean()
df['hours-per-week'].fillna(mean_hours,inplace=True)
```

### 1.2.6 Self-Check

Run the cell below to test the correctness of your code above before submitting for grading. Do not add code or delete code in the cell.

```
[15]: # Run this self-test cell to check your code;
# do not add code or delete code in this cell
from jn import testFillNa

try:
    p, err = testFillNa(df)
    print(err)
except Exception as e:
    print("Error!\n" + str(e))
```

Correct!

Check if we successfully converted all missing values to the mean value. Display the sum of missing values for the age column.

```
[16]: np.sum(df['age'].isnull(), axis = 0)
```

```
[16]: 0
```

In the code cell below, do the same for the hours-per-week column. Save the result to variable `sum_hours`.

### 1.2.7 Graded Cell

The cell below will be graded. Remove the line "raise NotImplementedError()" before writing your code.

```
[23]: # YOUR CODE HERE
sum_hours=np.sum(df['hours-per-week'].isnull(),axis=0)
```

### 1.2.8 Self-Check

Run the cell below to test the correctness of your code above before submitting for grading. Do not add code or delete code in the cell.

```
[24]: # Run this self-test cell to check your code;
# do not add code or delete code in this cell
from jn import testSumHours

try:
    p, err = testSumHours(df, sum_hours)
    print(err)
except Exception as e:
    print("Error!\n" + str(e))
```

Correct!

```
[ ]:
```